A Survey of SNARKs in Decentralized Finance Systems

Isaac Thomas[1]

Computer Science & Engineering, UC San Diego
isthomas@ucsd.edu

Abstract. Succinct non-interactive arguments of knowledge (SNARKs) have seen increasing adoption in decentralized finance (DeFi) systems, where scalability and privacy preservation are of high concern in verifying computation performed on blockchain networks. In particular, SNARKs using pairing-based cryptography are widely used in blockchain infrastructure/applications due to their small proof sizes, fast verification, and easily endowed zero-knowledge properties. In this survey, we review advancements in pairing-based SNARKs relevant to the performance, flexibility and confidentiality of blockchain networks and their hosted applications. Along with fundamental SNARK properties and early work, we discuss alternate characterizations of the complexity class NP, cryptographic primitives, and optimizations modern pairing-based SNARKs collectively employ to verify payments and execution paths through on-chain programs; and we weigh concrete security/performance/privacy tradeoffs emerging from such design choices. Finally, we discuss applications and future SNARK research avenues relevant to verifying on-chain computation.

1 Preliminaries

1.1 What is a SNARK?

Loosely speaking, a succinct non-interactive argument of knowledge (SNARK) is a short, one-shot proof that one *knows* some "witness" or value w satisfying a claim C related to public input(s) x. In practice, C is a claim that $x \in L$, where $L \in$ NP. We say that a *prover* $\mathcal{P}$ would send a such a proof to a *verifier* $\mathcal{V}$, which checks the proof against the claim to be proven and either accepts the proof or rejects it. We give a more formal definition below.

A more formal definition Let the relation $\mathcal{R} = \{(x, w)\} \subseteq \{0, 1\}^* \times \{0, 1\}^*$ be decidable in deterministic polynomial time, where $|w| = \mathsf{poly}(|x|) \; \forall (x, w) \in \mathcal{R}$. Let $\lambda \in \mathbb{N}$ denote the security parameter. Define the language $L = \{x \mid \exists w \text{ s.t. } (x, w) \in \mathcal{R}\}$. A succinct non-interactive argument of knowledge is the tuple of PPT algorithms (Setup, Prove, Verify) where:

- Setup$(1^\lambda, \mathcal{R}) \to \sigma$ receives an input security parameter and relation, and outputs a common reference string (CRS) σ containing the "setup" terms the prover and verifier will use to construct/verify proofs
- Prove$(\sigma, \mathsf{x}, \mathsf{w}) \to \pi$ receives the CRS, string x, and witness w and outputs a proof π
- Verify$(\sigma, \mathsf{x}, \pi) \to \{0, 1\}$ receives the CRS, string x, and proof π; it outputs 1 (accepts) or 0 (rejects)

and Setup, Prove, Verify have the following properties:

- **completeness:** If $\pi \leftarrow$ Prove$(\sigma, \mathsf{x}, \mathsf{w})$ and $(x, w) \in \mathcal{R}$ then Verify outputs 1 except with negligible probability.
- **knowledge soundness:** For all possible PPT algorithms Prove$'$, if $\pi' \leftarrow$ Prove$'(\sigma, \mathsf{x})$ and Verify$(\sigma, \mathsf{x}, \pi')$ outputs 1 with non-negligible probability, then there exists a PPT extractor $\mathcal{E}$ which, given access to the same inputs as Prove and $\mathcal{P}$'s random tape, outputs w s.t. $(x, w) \in \mathcal{R}$. In other words, if the proof is accepted then with high probability the prover must have known & used a valid witness to generate the proof.
- **succinctness:** Let $(x, w) \in \mathcal{R}$ and $\pi \leftarrow$ Prove$(\sigma, \mathsf{x}, \mathsf{w})$. Setup complexity and prover complexity are quasilinear in $|x|$, $|\pi|$ is sublinear in $|x|$, and verifier complexity is sublinear in $|x|$.
- **non-interactivity:** π is generated without any interaction between $\mathcal{P}$ and $\mathcal{V}$ after setup.

As an example, L could be the set of string encodings of satisfiable boolean circuits. The claim to be proven is that $x \in L$, where x is the encoding of the circuit in question; since L is fixed here, the claim can just be represented by x. We can then call x the *instance* or *statement*. Our witness w is our potentially secret "certificate" proving the validity of the claim; one could generate w by "executing" the circuit on the satisfying input, recording all the values the wires take on in the process. In a more realistic setting, L could be the set of tuples (P, x, y, z) corresponding to on-chain programs P producing the output z when executed on inputs x, y. Then our witness w could contain x, y and a trace of P s execution, while our statement ("the output of P is z") would be (P, z).

1.2   Early work

SNARKs are preceded by foundational work in complexity theory, probabilistic proofs, and cryptography. First was the formalization of interactive proofs by Goldwasser et al. [GMR19], which established a theoretical framework by which a prover $\mathcal{P}$ exchanges a finite number of messages during *interaction* with a verifier $\mathcal{V}$, which issues *random challenges* to $\mathcal{P}$ in order to verify their claim. They also introduced the notion of zero-knowledge via which a proof system leaks no other information beyond the validity of the claim being verified. Arora et al. [AS98, ALM$^+$98] later drew the seminal equivalence between NP and PCP$[O(\log n), O(1)]$, yielding a powerful characterization of NP as the set of languages with verifiers needing logarithmic randomness bits and *constant queries*, where n is the statement size. Given results showing all NP languages can have *very* efficient verifiers, the following question naturally arises: "short enough" or *succinct* arguments of knowledge exist for NP? Kilian's construction [Kil92] realized such a thing using merkle tree commitments to probabilistically checkable proofs (PCPs), and Micali later showed how to make his argument non-interactive [Mic94] using the Fiat-Shamir heuristic [FS87]. Despite these breakthroughs, SNARKs constructed from PCPs were undesirable due to massive overhead incurred in proof construction. Many approaches instead use algebraic characterizations of NP which admit means to efficient polynomial equality/divisibility testing. Additionally integrating "compiled" interaction and suitable cryptographic commitment schemes brings us closer to the practical SNARK schemes used in real-world blockchain systems.

As seen in Kilian's work, along this path to practically succinct argument systems lies an interesting avenue: considering only computationally bounded adversaries. In doing so, one can now use cryptographic schemes to produce small encrypt proof elements and quickly verify their integrity, while ensuring the intractability of breaking soundness. Pairing-based cryptography, which allows one to check algebraic relationships between quantities "in the exponent" using bilinear functions (behaving like multiplication) on elliptic curve group elements, spurred a sequence of works leading to methods practical enough for real-world blockchain applications. Though faster, all of these methods require stronger security assumptions. In particular, pairing-based methods derive their security in part from the elliptic curve discrete log assumption, which is stronger than the collision resistance assumptions used in other hashing-based methods [BSBHR19, AHIV17, WTS$^+$18, BSCR$^+$19]. Within and across denominations of SNARKs, such tradeoffs between performance, security, and flexibility are recurring points of comparison; and where they appear becomes clearer with a picture of which components are stitched together to form a SNARK.

1.3   A general "workflow" for SNARKs

Fig. 1. A general workflow describing how SNARKs reduce checking computation to checking polynomials. The highlighted path shows the course pairing-based approaches take. Acronyms: AET = algebraic execution trace, AIR = algebraic intermediate representation, CC = cryptographic compilation, PCW = polynomial codeword, PEQ = polynomial equality, PC = pairing check, PPT = polynomial proximity testing, R1CS = rank-1 constraint system.

Modern SNARKs largely use the "workflow" depicted in figure 1 to prove/verify computation. Here there are two courses; one reduces to checking polynomial relationships using pairing-based cryptography (left), and the other reduces to checking proximity to a Reed-Solomon code via combination of coding theoretic results and merkle trees. In an interactive setting, the "checks" or "testing" steps might involve randomized interaction with the verifier; but many interactive-proof-based methods use Fiat-Shamir to have the prover simulate such interaction, which the verifier can check. To balance coverage, depth, and accessibility, this paper focuses on approaches involving pairing-based cryptography (highlighted paths). We leave the surveying of methods using hashing-based cryptography as future work.

1.4 Representing the complexity class NP

Crucial to modern SNARKs is their representation of NP. How the computation to verify is "framed" directly influences how synergistic each method is with algebraic objects like polynomials and the cryptographic methods applied to them. This in turn affects the proof size and prover/verifier complexity, all of which are highly relevant to blockchain infrastructure/applications using SNARKs. Clearly, one must use an NP-complete language compatible with other aspects of the method in question. The most immediate choice of representation is boolean circuit satisfiability; in light of the use of polynomials, a reasonable sequel is *arithmetic* circuit satisfiability, in which the circuit of importance uses addition gates in place of OR gates, and multiplication gates in place of AND gates. Wires can also take on values in some prime field F_p rather than just 0 or 1. A natural question arises: are there representations of NP that seamlessly connect arithmetic circuit satisfiability with polynomial equality testing? The answer is yes, and we discuss notable examples of this below.

Rank 1 Constraint System (R1CS) Let $m, n \in \mathbb{N}$, and consider an arithmetic circuit with m gates and n input variables. with A rank-1 constraint system (R1CS) consists of a vector $w \in \mathbb{F}_p^n$, the matrices $L, R, O \in \mathbb{F}_p^{m \times n}$ and the relationship

$$Ow = (Lw) \circ (Rw) \tag{1}$$

A rank-1 constraint system (R1CS) *arithmetizes* the relationships between the left inputs, right inputs, and outputs of each gate, respectively. It assumes that the circuit has been preprocessed so that only multiplication gates remain, with the addition gates instead expressed as sums input to the multiplication gates. Intuitively, row i of L can be seen as a "selector" for the values in w that, when linearly combined via row L_i, form the left input to gate i. The same logic applies for R and O. Applying this logic in aggregate yields the above relationship between a Hadamard product of vectors and the desired output vector - in other words, three rank-1 matrices (hence "rank-1"). This problem is NP-complete with a relatively simple reduction from arithmetic circuit-SAT (follows easily from the loose definition above). One can also view it as a consolidation of the equations in the constraint system used by Bootle et al. [BCC+16], which uses two separate equations for multiplication gate constraints and linear constraints.

Quadratic Arithmetic Program (QAP) Connecting the R1CS representation of NP with polynomials involves converting either the rows or the columns of L, R, O into polynomials and creating some relationship that holds if and only if $Ow = (Lw) \circ (Rw)$. Given that for the i-th gate we have

$$(Oz)_i = \left(\sum_{j=1}^n L_{ij} w_j \right) \left(\sum_{j=1}^n R_{ij} w_j \right) \tag{2}$$

It could make sense to create a polynomial for the j-th witness variable which evaluate to each L_{ij} given the gate i. This would involve interpolating column j of the matrix L into a polynomial $u_j(x)$. The same intuition applies to the matrices R, O (using $v_j(x), w_j(x)$ respectively) and is the idea behind the quadratic arithmetic program (QAP) introduced by Gennaro et al. [GGPR13]. We focus on the definition of a regular QAP.

Definition 1 (regular QAP). *A regular Quadratic Arithmetic Program Q over a field $\mathbb{F}_p$ comprises:*

- *The target polynomial $t(x) = \prod_{i=1}^m (x - i)$, where m is the number of gates*
- *The polynomials $\{u_i(x), v_i(x), w_i(x)\}_{i=0}^n$, where $u_i(j) = L_{ji}$, $v_i(j) = R_{ji}$, and $w_i(j) = O_{ji}$*

  Q is satisfied by a witness $c \in \mathbb{F}_p^n$ if and only if:

$$p(x) = \left(\sum_{i=0}^n c_i u_i(x) \right) \cdot \left(\sum_{i=0}^n c_i v_i(x) \right) - \left(\sum_{i=0}^n c_i w_i(x) \right) = h(x)t(x) \equiv 0 \mod t(x) \tag{3}$$

Here $t(x)|p(x)$ is synonymous with $p(x)$ vanishing at all points which are gate identifiers, which will be true if the gate relationship is indeed satisfied by the given inputs, and will not be true otherwise. QAP satisfiability is NP-complete, as can be shown via reduction from R1CS-SAT or arithmetic circuit-SAT.

Though not exhaustive by any means, these two representations of NP yield a connection between the computation being verified, arithmetic circuit satisfiability, and an instance of polynomial divisibility testing. From this point, the polynomial divisibility check can be verified in a hidden fashion using pairing-based cryptography. This idea is used by virtually all pairing-based methods, with modifications to the constraint system and the polynomial relationship being checked.

1.5   Polynomial Properties

As mentioned, modern SNARKs reduce proving/verifying statements about computation to polynomial identity/property testing via suitable characterizations of NP. The use of polynomials over other mathematical objects is justified by how little information one needs to distinguish between two polynomials, which is a consequence of the following lemma.

Lemma 2. *Schwarz-Zippel lemma Let $f(x_1, x_2, \ldots x_n) \in R[x_1 \ldots x_n]$ be a nonzero polynomial in n variables defined over an integral domain $\mathbb{F}^n$. Suppose the element $(a_1, a_2, \ldots a_n)$ is selected uniformly at random from a finite subset $S \subset \mathbb{F}^n$. Then*

$$\Pr(f(a_1, a_2, \ldots a_n) = 0) \leq \frac{deg(f)}{|S|}$$

where $deg(f)$ is the maximum sum of the degrees of any term's variables. It immediately follows that if $f = g - h$, then for a randomly sampled point $(a_1, \ldots a_n) \in \mathbb{F}$, we have that $\Pr(g(a_1, \ldots a_n) = h(a_1, \ldots a_n))$ is bounded from above by $\frac{d}{|\mathbb{F}|}$. In other words, g and h will output different values at $(a_1 \ldots a_n)$ with high probability if they are not equal, assuming the d and $\mathbb{F}$ are chosen so that $d \ll |\mathbb{F}|$. Thus, given the right choice of degree and field size, a verifier will still be able to distinguish between two unequal polynomials with high probability using a single evaluation point. Succinct argument systems can make use of this to check polynomial equalities at a single point with small proof size, fast verification, and little probability of spurious equality.

1.6   Pairing-based Cryptography

In both its conception and current application, pairing-based cryptography [BF01] allows one to check algebraic relationships between quantities "in the exponent" using bilinear functions on arbitrary source group elements (usually from elliptic curve groups) as this works well with checking polynomial equalities, it is frequently used in modern SNARKs. Elliptic curve groups are preferred here due to small key sizes – determined by the size of the base field the point coordinates are from. There is also a group operation over elliptic curve points behaving like addition, in which one computes a sum of points $A + B$ by reflecting over the x-axis the point on both the curve and the line AB (if $A = B$ we use the tangent line through A). We can then define multiplication of a point A by some $s \in \mathbb{F}_s$ in the curve's *scalar field* as summing s copies of A. Here, s can be considered an *elliptic curve discrete logarithm* and is believed to be hard to compute given A and sA; this is the elliptic curve variant of the discrete log assumption. The hardness of this problem yields a scheme for additively homomorphic encryption: given a cyclic multiplicative group $\mathbb{G}$ over some wisely-chosen elliptic curve and its generator G, encrypt x as $xG \in \mathbb{G}$. Pairings can render this scheme *doubly* homomorphic due to the properties of bilinear pairings, which behave like multiplication (linear in each input). This paper abstracts elliptic curve groups and pairings for brevity; we instead focus on the properties they exhibit which make pairing-based cryptography useful in this setting.

More formally, suppose we have two cyclic elliptic curve groups $\mathbb{G}_1$ and $\mathbb{G}_2$ generated by elements G_1 and G_2, respectively. We denote the scalar multiple of a point $P \in \mathbb{G}_i$ by kP, which is just the same as k additions of P and $k \in \mathbb{F}_s$ where $\mathbb{F}_s$ is some scalar field. For a bilinear pairing $e : \mathbb{G}_1 \times \mathbb{G}_2 \to \mathbb{G}_t$ where $P_1 \in \mathbb{G}_1$ and $P_2 \in \mathbb{G}_2$, we have the following properties:

- for $a, b \in \mathbb{F}$, $e(aP_1, bP_2)e(P_1, P_2)^{ab} = e(P_1, P_2)^{ba} = e(bP_1, aP_2)$
- for $Q_1 \in \mathbb{G}_1$, $e(P_1 + Q_1, P_2) = e(P_1, P_2)e(Q_1, P_2)$
- for $Q_2 \in \mathbb{G}_2$, $e(P_1, P_2 + Q_2) = e(P_1, P_2)e(P_1, Q_2)$
- $e(G_1, G_2)$ generates $\mathbb{G}_t$ (non-degeneracy)

These properties make it easy to check multiplicative relationships. For instance, suppose we have three polynomials $A(x), B(x), C(x)$ and we want to check that $A(\alpha)B(\alpha) = C(\alpha)$ for some input α. From bilinearity of e it follows that

$$e(G_1, G_2)^{A(\alpha)B(\alpha)} = e(G_1, G_2)^{C(\alpha)}$$
$$\Leftrightarrow e([A(\alpha)]_1, [B(\alpha)]_2) = e([C(\alpha)]_1, G_2)$$

where $[A(\alpha)]_1 = A(\alpha)G_1$ (same idea for the other values). So if we can reduce checking circuit satisfiability to checking a polynomial relationship, we can send elliptic curve points as proof elements with which the verifier would perform such a pairing check. The complexity of a naive pairing computation is quasi-quadratic in the target group size, but there are pairing-friendly choices of elliptic curve groups which make this operation faster.

1.7   Polynomial commitment schemes (PCS)

A polynomial commitment scheme is a protocol by which a prover claims they know a polynomial $f(x)$ satisfying some relationship, and a verifier checks this claim. A noteworthy example of such a relationship is that $f(y) = z$ for some fixed y, z. We detail a noteworthy polynomial commitment scheme for this exact task known as the Kate-Zaverucha-Goldberg commitment scheme (KZG) [KZG10]. Suppose P wants to prove they know f of degree d such that $f(\beta) = z$. It follows that $f(X) - z$ is divisible by $(X - \beta)$, so we should be able to construct $h(X) = \frac{f(X) - z}{X - \beta}$ since we know f. P commits to f and h by their hidden evaluations on some α agreed upon in advance by the P and the verifier V. In practice, the evaluations are hidden via scalar multiplication by an elliptic curve group generator $g_1 \in \mathbb{G}_1$. This scheme uses two elliptic curve groups for this purpose, so assume the point g_2 generates the EC group $\mathbb{G}_2$. Let $[a]_i = ag_i \in \mathbb{G}_i$. Suppose $\mathcal{P}$ and $\mathcal{V}$ agree on a *trusted setup* containing the hidden terms $\{[1]_1, [\alpha]_1, [\alpha^2]_1, \ldots [\alpha^d]_1, [1]_2, [\alpha]_2\}$. P sends $[f(\alpha)]_1$ and $[h(\alpha)]_2$. V then sends a challenge point γ to which $\mathcal{P}$ responds with $[f(\gamma)]_1$ and $[h(\gamma)]_1$. V then checks that $f(\gamma) - z = h(\gamma)(\gamma - \alpha)$ where h was constructed from the evaluation requirement we wanted to prove that f satisfies. For a bilinear pairing function $e : \mathbb{G}_1 \times \mathbb{G}_2 \to \mathbb{G}_t$, $\mathcal{V}$ checks that

$$e([f(\gamma) - z]_1, g_2) = e([h(\gamma)]_1, [\gamma]_2 - [\alpha]_2) \tag{4}$$

$$\Leftrightarrow e(g_1, g_2)^{f(\gamma) - z} = e(g_1, g_2)^{h(\gamma)(\gamma - \alpha)} \tag{5}$$

$$\Leftrightarrow f(\gamma) - z = h(\gamma)(\gamma - \alpha) \tag{6}$$

KZG commitment schemes require constant size communication since the polynomials involved can be collapsed to a point without losing virtually any distinguishability. Given their use of elliptic curve groups to meet this end, their security depends on the hardness of computing elliptic curve discrete logarithms (find α given the points g and αg), for which no classical polynomial time algorithm is known [Die11]. The drawbacks of using this scheme include the need to generate a trusted setup and the lack of post-quantum security. Being quite customizable, KZG works with batching, multivariate polynomials, and various optimizations for faster polynomial operations. That said, many methods use different polynomial commitment schemes, some of which avoid pairing-based cryptography entirely. Nonetheless we opt for this example to give an idea of commitment schemes making use of pairings, as this is relevant to many modern pairing-based SNARKs.

2   Circuit-specific pairing-based SNARKs

The seminal PCP theorem results [AS98, ALM+98] yielded a powerful characterization of NP as a set of languages with polynomial-time PCP verifiers only needing to view a constant number of proof bits *regardless of proof size*. However, early work on pairing-based SNARKs by Gennaro et al. [GGPR13] was motivated by the possibility of succinct arguments of knowledge using a representation *more suitable* than PCPs for integration with cryptographic primitives. Several other works [**WHICH ONES**] came close to meeting this end but could not attain sublinear proof size.

To this end, Gennaro et al. [GGPR13] coined quadratic span programs (QSP) and quadratic arithmetic programs (QAP). These representations of NP allow one to reduce checking a computation to checking a multiplicative relationship between two polynomials in a way that "computes" the circuit representing the computation. Although the QSP construction is noteworthy, we limit discussion to their QAP-related construction due to the clearer connection to arithmetic circuits and downstream polynomial checking. We recall the strong QAP form here:

$$\underbrace{\left(\sum_{j=1}^{n} a_i v_i(x)\right)}_{\text{left inputs}} \cdot \underbrace{\left(\sum_{j=1}^{n} b_i w_i(x)\right)}_{\text{right inputs}} - \underbrace{\left(\sum_{j=1}^{n} c_i y_i(x)\right)}_{\text{outputs}} = h(x)t(x) \equiv 0 \mod t(x) \tag{7}$$

This representation is compiled into cryptographic proofe elements via an additively homomorphic encoding E – namely, a one-to-one, function for which addition operations are preserved in the output space, and inversion is difficult. Given an instance of arithmetic circuit satisfiability one is trying to prove/verify, E serves to compile the corresponding QSP/QAP instance checking polynomial divisibility (by $t(x)$) into a proof containing just 9 group elements (non-ZK). We detail the non-zero-knowledge version here for clarity, denoting $E(x)$ by $[x]$. The prover $\mathcal{P}$ computes the quotient polynomial $h(x)$ and sends

$$\pi = \Big([v_{mid}(s)], [w(s)], [y(s)], [h(s)], \tag{8}$$

$$[\alpha v_{mid}(s)], [\alpha w(s)], [\alpha y(s)], [\alpha h(s)], \tag{9}$$

$$[\beta_v v_{mid}(s) + \beta_w w(s) + \beta_y y(s)] \Big) \tag{10}$$

where $\alpha, s, \beta_v, \beta_w, \beta_y \in \mathbb{F}_p$, are secret random preprocessing elements; $v_{mid}(s), w(s), y(s)$ are the witness-weighted combinations of the wiring polynomials in the QAP; and $h(s)$ is the evaluation of the quotient polynomial an honest prover would know. In particular, $v_{mid}(x) = \sum_{k \in \mathcal{I}_{mid}} a_k v_k(x)$ where $\mathcal{I}_{mid}$ are the indices corresponding to private circuit inputs. $w(x), y(x)$ are definedn similarly, but over all $k \in \{0 \dots m\}$. The verifier in turn checks 5 equations using the additive properties of E. In particular, the verifier receives the following proof π' *which it does not yet know is valid*, hence the new notation of elements reminiscent of those in π:

$$\pi' = (\pi_{v_{mid}}, \pi_w, \pi_y, \pi_h, \pi_{v'_{mid}}, \pi_{w'}, \pi_{y'}, \pi_{h'}, \pi_{z'}) \tag{11}$$

$\mathcal{V}$ then checks 6 pairing equations involving the components of π' (recall the pairing notation from earlier section):

$$e([v_0(s) + v_{in}(s) + v_{mid}(s)]_1, [w_0(s) + w_{in}(s) + w_{mid}(s)]_2) = \tag{12}$$

$$e([t(s)]_1, [h(s)]_2) \cdot e([y_0(s) + y_{in}(s) + y_{mid}(s)]_1, G_2) \tag{13}$$

$$e([v'_{mid}(s)]_1, G_2) = e([v_{mid}(s)]_1, [\alpha]_2) \tag{14}$$

$$e([w'(s)]_1, G_2) = e([w(s)]_1, [\alpha]_2) \tag{15}$$

$$e([y'(s)]_1, G_2) = e([y(s)]_1, [\alpha]_2) \tag{16}$$

$$e([h'(s)]_1, G_2) = e([h(s)]_1, [\alpha]_2) \tag{17}$$

$$e([z'(s)]_1, [\gamma]_2) = e([\beta_v v_{mid}(s) + \beta_w w(s) + \beta_y y(s)]_1, [\gamma]_2) \tag{18}$$

where $v'_{mid}(s)$ is encoded in $\pi_{v'_{mid}}$, $w'(s)$ is encoded in $\pi_{w'}$, and so on. The first one checks the QAP divisibility relation. The next four equations are based on knowledge-of-exponent assumptions; they check if $\mathcal{P}$ knows the wiring polynomials used in the proof. The last equation checks that $\mathcal{P}$'s proof used wiring polynomials consistent with those agreed upon in the trusted setup.

GGPR uses a *strong* QAP for their scheme, which uses a different set of coefficients for each sum in the QAP equation. This mandates a *strengthening step* in which extra constraints between the coefficient sets are materialized, tripling prover work and preprocessing size. Parno et al. [PHGR16] made the simple optimization of using a *regular* QAP, which uses the same set of coefficients a_i for each sum in the QAP expression, contrary to the a_i, b_i, c_i seen in the strong QAP form. This change eliminated the need for the strengthening step without any noteworthy compromises, although it required modifications to proof elements and verification checks that ensure consistent use of coefficients in QAP terms. Another improvement was the use of an asymmetric pairing operation $e : \mathbb{G}_1 \times \mathbb{G}_2 \to \mathbb{G}_t$ such that $\mathbb{G}_1 \neq \mathbb{G}_2$, which in practice is 3-4x as efficient as its symmetric counterpart ($\mathbb{G}_1 = \mathbb{G}_2$). Imaginably, both of these optimizations appear in following QAP-based approaches, including the current state of the art.

With these two methods as a springboard, Groth [Gro16] used stronger security assumptions to make massive improvements in proof size, setup size, and verifier complexity. Groth still used the regular QAP representation; but unlike previous methods that compiled each polynomial evaluation into two group elements, he compiled the wiring polynomials to just one group element (ex. just $w(s)$ instead of $w(s), \alpha w(s)$). This eliminates multiple setup terms, reduced proof size to 3 group elements, and reduced verifier operations to a single pairing check using three pairings. This imaginably resulted in a massive performance increase across multiple fronts over previous methods. However, this single-element compilation prevents use of knowledge-of-exponent assumptions, so security is argued in the generic group model where, loosely speaking, adversaries cannot exploit any group specific properties. While seemingly strong, such assumptions are not necessarily uncalled for, as Gentry et al. have already shown that SNARKs for NP cannot exist without non-falsifiable assumptions to begin with [GW11]. Reduction in proof size and verifier complexity could also be regarded as "worth the security assumptions" on this front. Consequently, Groth's work (colloquially known as groth16) saw large adoption in blockchain systems verifying homogeneous computation; applications like Tornado Cash [PSS19] use this method to verify withdrawals from privacy-preserving pools without exposing source-destination address relationships; and UTXO-based privacy-preserving blockchains like Zerocash network [BSCG+14] used groth16 to prove/verify presence of unspent funds in the merkle tree inherent to the UTXO model.

3   Circuit-independent (universal) pairing-based SNARKs

Homogeneous computation aside, circuit-specific pairing-based methods have a couple of noteworthy shortcomings. The first of these, as noted by Groth [GKM+18], is that there is no way of ensuring the deployers of the SNARK in question have disposed of the secret randomness used to generate the trusted setup. The second and more important issue is a lack of flexibility due to the use of computation-dependent (non-universal)

Table 1. Comparison of Pairing-based SNARK Systems

method	SRS size	prover-time	proof-length	verifier-time	universal	updatable	assumptions
GGPR13	$O(n)$	$O(n \log n)$	$7\mathbb{G}_1 + 1\mathbb{G}_2 + 1\mathbb{F}$	$O(n)$	No	No	q-PKE, q-PDH
PGHR13	$O(n)$	$O(n \log n)$	$7\mathbb{G}_1 + 1\mathbb{G}_2$	$O(n)$	No	No	q-PDH
Groth16	$O(n)$	$O(n \log n)$	$2\mathbb{G}_1 + 1\mathbb{G}_2$	$O(1)$	No	No	q-type
GMKL18	$O(n^2)$	$O(n \log n)$	$O(1)\mathbb{G}_1 + O(1)\mathbb{G}_2 + O(1)\mathbb{F}$	$O(1)$	Yes	Yes	SXDH
MBKM19	$O(n)$	$O(n \log n)$	$7\mathbb{G}_1 + 1\mathbb{G}_2 + 5\mathbb{F}$	$O(1)$	Yes	Yes	SXDH
GWC19	$O(n)$	$O(n \log n)$	$7\mathbb{G}_1 + 1\mathbb{G}_2$	$O(1)$	Yes	Yes	SXDH
CHM+19	$O(n)$	$O(n \log n)$	$7\mathbb{G}_1 + 1\mathbb{G}_2 + 6\mathbb{F}$	$O(1)$	Yes	Yes	AGM

preprocessing. Consider the groth16 trusted setup for instance:

$$\sigma = \Big([\alpha]_1, [\beta]_1, [\gamma]_1, [\delta]_1, \{[x^i]_1\}_{i \in [0, n-1]}, \{[\gamma^{-1}(\beta u_i(x) + \alpha v_i(x) + w_i(x))]_1\}_{i \in [0, \ell]}, \tag{19}$$

$$\{[\delta^{-1}(\beta u_i(x) + \alpha v_i(x) + w_i(x))]_1\}_{i \in [\ell+1, m]}, \tag{20}$$

$$\{[\delta^{-1} x^i t(x)]_1\}_{i \in [0, \deg(h)]}, [\beta]_2, [\gamma]_2, [\delta]_2, \{[x^i]_2\}_{i \in [0, n-1]} \Big) \tag{21}$$

As is the case with other discussed methods, the contributions of the i-th witness component to gate g are "stuck" in the encodings of $u_i(x), v_i(x), w_i(x)$ (not the component's value, but its index). Thus if the circuit structure changes, the setup terms encoding the hidden combinations of the $v_i(x), w_i(x), y_i(x)$, etc. must be regenerated. A natural question arises: why not just update the setup somehow if the computation structure changes? This is not possible for two reasons, either of which is sufficient to deter updatability. The first is that, as mentioned the hidden $u_i(x), v_i(x), w_i(x)$ are interpolated from relationships between witness values (secret) and the gates they feed into. Adding a new gate and its associated wire connections would require adding a new data point to each of these polynomials. Thus, one would have to re-interpolate all of them which is not possible without violating cryptographic assumptions (since they are encrypted) or regenerating the entire setup, which defeats the purpose of updating anything. The other reason concerns the setup's combination of these hidden polynomial terms, such as the $[\beta u_i(x) + \alpha v_i(x) + w_i(x)]_1$. Groth [GKM$^+$18] showed these setup terms could be used to extract the the constituent monomials and ultimately break soundness if the setup allowed updates. These observations suggest that the conception of a QAP-based updatable SNARK is unlikely, and that different approaches are necessary. Such approaches should be *universal*, in the sense that they allow proving computation of any structure up to a certain size; and *updatable*, meaning that the trusted setup can be efficiently updated by any party and remains sound if at least one setup contributor is honest. This idea becomes highly relevant when deploying blockchain systems using SNARKs with trusted setups; users would benefit from knowing the deployers cannot subvert the network, and that anyone can perform security-preserving updates to the SNARK system being used.

To address these shortcomings, Groth [GKM$^+$18] produced a multivariate scheme that encoded QAP elements differently and achieved constant proof size and verification complexity. Although this scheme involved a linear-time procedure by which a linear-size circuit-specific setup could be produced as needed, true universality required quadratically many terms w.r.t circuit size. Furthermore, SRS updates take a quadratic number of group exponentiations and update verification a linear number of pairing operations. For circuits with millions of gates this is impractical. However, this hints that a SNARK with linear-size universal & efficiently updatable setup could be attainable if the setup terms are univariate (and monomial). Maller et al. achieved this with Sonic, which draws inspiration from techniques of Bootle et al. [BCC$^+$16] reducing circuit satisfiability to checking Laurent polynomials encoding a Hadamard matrix product (models mul. gate operations) paired with a linear constraint system (models wire connections & addition gates). Though they take up more space due to the presence of X^{-i} term for every X^i term (loosely speaking), the setup contains only univariate monomials and therefore requires linear space. Being composed of monomials, the setup can be updated by any party, for which they supply a proof of correctness. In a realistic deployment of this scheme, the setup would not be circuit dependent, and a single update could eliminate the risk that the deployers still hold valid toxic waste that can be used to subvert the setup.

Though significant in its own right, Sonic suffers from large constants in proof construction complexity despite being asymptotically quasilinear, with Sonic more negatively affected in this regard. A likely cause is the attempt to accommodate for n-fan-in circuits whose gates can accept arbitrarily many inputs. While a reasonable generalization, this allows any given linear constraint the ability to use arbitrarily many witness inputs, requiring the use of entire witness and selector vectors for each constraint. It also causes a bloated

permutation argument since a given gate may require arbitrarily many copy constraints between its own inputs and outputs from preceding gates that feed into it. Gabizon et al. [GWC19] addressed these issues with a simplifying assumption that may not catch one's eye, but turns out to have important performance implications: simply let each circuit gate take two inputs. The impact of this is twofold. Firstly, it enables a much simpler constraint representation of arithmetic circuit-SAT, where for an n gate circuit we have the following constraint for gate i. Here a_i, b_i, c_i the indices in $\mathbf{x}$ corresponding to the left input, right input, and output of the i-th gate, respectively; and $\mathbf{q_L}, \mathbf{q_R}, \mathbf{q_O}, \mathbf{q_M}, \mathbf{q_C}$, are "selector vectors" determining which gate-related values partake in the constraint. This allows one to express addition or multiplication gates in the same constraint by setting $(\mathbf{q_L})_i, (\mathbf{q_R})_i, (\mathbf{q_M})_i$ accordingly.

$$(\mathbf{q_L})_i \cdot \mathbf{x}_{(\mathbf{a})_i} + (\mathbf{q_R})_i \cdot \mathbf{x}_{(\mathbf{b})_i} + (\mathbf{q_O})_i \cdot \mathbf{x}_{(\mathbf{c})_i} + (\mathbf{q_M})_i \cdot (\mathbf{x}_{(\mathbf{a})_i} \cdot \mathbf{x}_{(\mathbf{b})_i}) + (\mathbf{q_C})_i = 0 \tag{22}$$

During proof construction, these terms are collected into polynomials $q_L(x), a(x), q_R(x), b(X)$, etc. and combined in the same manner as the original constraint equation. Secondly, the 2-fan-in assumption enables optimizations to the permutation argument inspired by Groth et al. and Maller et al. [BG12, MBKM19]. While sonic used a product check accounting for the arbitrarily many copy constraints per gate, each product check per gate involves three terms in the numerator and denominator each, resulting in a "grand product" used in proof construction. We show the non-zero-knowledge version here, but adding ZK properties requires shifting this equation by adding a "masking" polynomial that still vanishes on the desired domain in order to not corrupt the expression.

$$z(X) = L_1(X)+ \tag{23}$$

$$\sum_{i=1}^{n-1} L_{i+1}(X) \prod_{j=1}^{i} \frac{(w_j + \beta\omega^j + \gamma)(w_{n+j} + \beta k_1 \omega^j + \gamma)(w_{2n+j} + \beta k_2 \omega^j + \gamma)}{(w_j + \beta\sigma^*(j) + \gamma)(w_{n+j} + \beta\sigma^*(n+j) + \gamma)(w_{2n+j} + \beta\sigma^*(2n+j) + \gamma)} \tag{24}$$

where $L_i(x) = \frac{X^n - 1}{X - \omega^i}$ is the i-th Lagrange basis polynomial defined over the n-th roots of unity (gives 1 when $X = \omega^i$ and 0 otherwise). Notably, the polynomial check works well with a common optimization to PlonK's invocation of KZG scheme; by representing the hidden monomials in the Lagrange basis, the point evaluations of polynomials can be computed in $O(d)$ operations where d is the degree of the polynomial in question, which is faster than the $O(d \log d)$ complexity for interpolation via inverse FFT & subsequent evaluation. The pairing check performed by the verifiere is also "fixed argument" in the second source group, enabling both of only 2 pairings the verifier performs to be 30% faster than the traditional pairing [CS10].

The combination of flexibility, intuitive arithmetization, and better efficiency has made PlonK a gateway to other methods increasing expressiveness and performance, with important implications for so-called "zero-knowledge virtual machines" (zkVMs). Ambrona et al. [GW20b] proposed methods to optimize the constraint system used, as well as optimized circuits for PlonK-based verifiable implementations of the "SNARK-friendly" Poseidon hash function [GKR+21]. For zkVM operations less compatible with SNARK systems, Gabizon et al. proposed Plookup [GW20a] which adapts the PlonK permutation argument to verify that a set of values is present in some predetermined table; this is particularly useful for constraining the correctness of zkVM operations involving nonlinear operations that would make the resulting circuits / constraint system inefficient to verify. A notable example is a "SNARK-unfriendly" hash function like SHA-3 [D+15], which has many nonlinear bit-mixing operations. Variations like halo2 [Dev24] and plonky2 [POL22] combine the PlonK arithmetization with commitment schemes inspired by Bulletproofs [BBB+18] and FRI [BSBHR18] respectively to shed the need for a trusted setup at the expense of slower verification. These two advancements highlight the customizability of the PlonK system.

4 Applications

4.1 Privacy-preserving pools

Traditional blockchain network innerworkings expose transaction details, making privacy preservation non-trivial at the application layer. Fortunately, modern SNARKs can help address this issue when endowed with zero-knowledge properties. Applications like Tornado Cash [PSS19] allow one to deposit funds from wallet A, then withdraw the funds to a different wallet B without revealing the link between the two addresses on-chain. The user has to prove that they are the owner of the depositing wallet using a groth16 proof of deposit note membership in a Merkle tree maintained by the Tornado Cash smart contracts on-chain. The homogeneous omputation structure and need for succinct proofs/verification due to tight EVM smart contract gas limits make groth16 a suitable choice.

Fig. 2. Merkle tree with authentication path (in red) for commitment C_6 (in gray)

4.2 Privacy-preserving blockchains

Beneath the application layer there are blockchains utilizing zero-knowledge properties of SNARKs for infrastructural privacy preservation. For each transaction submitted, such systems must obscure the source, relevant currencies, amounts used, and payment recipients. Zerocash [BSCG+14] is a notable example of a network that does such a thing. Intended to be a privacy-preserving version of Bitcoin [Nak08], Somewhat similar to (and preceding) Tornado Cash, Zerocash the unspent-transaction-output (UTXO) model, where transactions create so-called "unspent funds" that the party who "owns" them is entitled to spend. When they do spend those funds, another record of "unspent funds" for the recipient of the spending is created. These UTXO records are stored in a merkle tree, and spending funds requires proving knowledge of a leaf in the tree corresponding to the funds. Early versions of Zerocash used groth16 for this, but the NU5 upgrade of 2022 switched to the universal halo2 proof system [Dev24], which enables a combination of a PlonK-style arithmetization and Bulletproofs-style commitment scheme [BBB+18] requiring no trusted setup at the expense of verifier performance.

4.3 Verifiable Virtual Machines (VVM)

Multiple "Zero-knowledge virtual machines" (zkVMs) [Scr24, Pol24, Suc23, RIS23, OO25] have emerged as architecture-specific and general solutions for proving validity of layer-2 transaction batches on the layer 1 (L1) chains settling them. We use the more accurate phrase "verifiable virtual machine" (VVM) since succinctness is a higher priority than zero-knowledge in most L2 architectures. Sitting at the execution layer of the blockchain node implementation, the VVM is responsible for executing calls to smart contracts, generating corresponding execution traces/witnesses, and generating proofs of correct execution for each batch of transactions. PlonK-style variants have been a popular first approach due to the diverse computation structure and lookup-style arguments for verifying "SNARK-unfriendly" operations. Many VVM implementations use circuit domain-specific languages (DSLs) like halo2 to implement the circuits constraining the computation details.

5 Future Research

5.1 Post-quantum algebraically-friendly SNARKs

Though multiple of the SNARKs discussed are practical from a performance standpoint, quantum computers may render them useless from a security standpoint in the near future. As a result, there is a strong effort to develop efficient enough post-quantum SNARKs. To this end, many works [BSBHR19, AHIV17, COS20, Set20, AST24] use hashing-based based Reed-Solomon proximity testing methods [BSBHR18], which relies on the hardness of computing hash function collisions. While such approaches are believed to be post-quantum secure, the cryptographic primitives used offer no additively homomorphic structure enabling the recursive composition of smaller proof elements & faster verification we see in pairing-based methods. Albrecht et al. [ACL+22] made some progress in this direction via a lattice-based SNARK with logarithmic-time verification; it derives security from the hardness of the short integer solution (SIS) problem in lattices. Future work that quantifies the security assumptions in this area and decreases prover/verifier complexity could realize practical SNARKs using additively homomorphic post-quantum encryption, which the blockchain ecosystem would eventually accept given sufficient tooling and support.

5.2 Automated verification of SNARK methods/applications

Currently, several layer 2 blockchains collectively hold tens of billions of dollars **[CITE]** protected by the integrity of SNARK methods used to verify block validity; if either the methods or their implementations

have soundness bugs, it could allow a knowledgeable adversary to prove an invalid state transition for financial gain, causing massive financial loss to users and the chain. There are two noteworthy and related directions of ongoing work in this domain. The first is related to formally verifying the SNARK methods themselves in various mechanized cryptographic models using Lean 4 [ark24, dMU] and other techniques for lookup-based arguments, as done with JOLT zkVM [AST24, KDT24]. Though very promising and easily updatable, there are still many gaps to fill. Furthermore, the periodic introduction of new methods will require constant updates to the relevant types and theorems. Nonetheless, seeking formal guarantees on SNARK behavior in various mechanized theoretical models will make them increasingly fit for use in large-scale DeFi infrastructure/applications needing to verify computation. A more adversarial direction concerns automated detection of missing circuit constraints [PCW$^+$23], which detect nondeterministic circuit behavior using symbolic execution and SMT solvers. Picus has been used heavily in verifying Risc Zero VM [RIS23] Preliminary work by Chaliasos et al. [CET$^+$24] has found that circuit-related soundness bugs are the most frequent kind of bug by a huge margin, suggesting that building upon this existing work will be important for the security of VVMs. For instance, it would be extremely useful to have tools detect missing constraints in circuits validating interacting sub-witnesses like in the STARK-based SP1 zkVM architecture, which uses LogUp [Hab22] to verify correctness of interactions.

References

ACL+22. M. R. Albrecht, V. Cini, R. W. Lai, G. Malavolta, and S. A. Thyagarajan. Lattice-based snarks: publicly verifiable, preprocessing, and recursively composable. In *Annual International Cryptology Conference*, pages 102–132. Springer, 2022.

AHIV17. S. Ames, C. Hazay, Y. Ishai, and M. Venkitasubramaniam. Ligero: Lightweight sublinear arguments without a trusted setup. In *Proceedings of the 2017 acm sigsac conference on computer and communications security*, pages 2087–2104, 2017.

ALM+98. S. Arora, C. Lund, R. Motwani, M. Sudan, and M. Szegedy. Proof verification and the hardness of approximation problems. *Journal of the ACM (JACM)*, 45(3):501–555, 1998.

ark24.   GitHub - Verified-zkEVM/ArkLib: Formally Verified Arguments of Knowledge in Lean — github.com. https://github.com/Verified-zkEVM/ArkLib, 2024. [Accessed 08-05-2025].

AS98.   S. Arora and S. Safra. Probabilistic checking of proofs: A new characterization of np. *Journal of the ACM (JACM)*, 45(1):70–122, 1998.

AST24.   A. Arun, S. Setty, and J. Thaler. Jolt: Snarks for virtual machines via lookups. In *Annual International Conference on the Theory and Applications of Cryptographic Techniques*, pages 3–33. Springer, 2024.

BBB+18. B. Bünz, J. Bootle, D. Boneh, A. Poelstra, P. Wuille, and G. Maxwell. Bulletproofs: Short proofs for confidential transactions and more. In *2018 IEEE symposium on security and privacy (SP)*, pages 315–334. IEEE, 2018.

BCC+16. J. Bootle, A. Cerulli, P. Chaidos, J. Groth, and C. Petit. Efficient zero-knowledge arguments for arithmetic circuits in the discrete log setting. In *Advances in Cryptology–EUROCRYPT 2016: 35th Annual International Conference on the Theory and Applications of Cryptographic Techniques, Vienna, Austria, May 8-12, 2016, Proceedings, Part II 35*, pages 327–357. Springer, 2016.

BF01.   D. Boneh and M. Franklin. Identity-based encryption from the weil pairing. In *Annual international cryptology conference*, pages 213–229. Springer, 2001.

BG12.   S. Bayer and J. Groth. Efficient zero-knowledge argument for correctness of a shuffle. In *Advances in Cryptology–EUROCRYPT 2012: 31st Annual International Conference on the Theory and Applications of Cryptographic Techniques, Cambridge, UK, April 15-19, 2012. Proceedings 31*, pages 263–280. Springer, 2012.

BSBHR18. E. Ben-Sasson, I. Bentov, Y. Horesh, and M. Riabzev. Fast reed-solomon interactive oracle proofs of proximity. In *45th international colloquium on automata, languages, and programming (icalp 2018)*. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik, 2018.

BSBHR19. E. Ben-Sasson, I. Bentov, Y. Horesh, and M. Riabzev. Scalable zero knowledge with no trusted setup. In *Advances in Cryptology–CRYPTO 2019: 39th Annual International Cryptology Conference, Santa Barbara, CA, USA, August 18–22, 2019, Proceedings, Part III 39*, pages 701–732. Springer, 2019.

BSCG+14. E. Ben Sasson, A. Chiesa, C. Garman, M. Green, I. Miers, E. Tromer, and M. Virza. Zerocash: Decentralized anonymous payments from bitcoin. In *2014 IEEE Symposium on Security and Privacy*, pages 459–474, 2014.

BSCR+19. E. Ben-Sasson, A. Chiesa, M. Riabzev, N. Spooner, M. Virza, and N. P. Ward. Aurora: Transparent succinct arguments for r1cs. In *Advances in Cryptology–EUROCRYPT 2019: 38th Annual International Conference on the Theory and Applications of Cryptographic Techniques, Darmstadt, Germany, May 19–23, 2019, Proceedings, Part I 38*, pages 103–128. Springer, 2019.

CET+24. S. Chaliasos, J. Ernstberger, D. Theodore, D. Wong, M. Jahanara, and B. Livshits. Sok: What don't we know? understanding security vulnerabilities in snarks, Feb 2024. arXiv:2402.15293v1.

COS20.   A. Chiesa, D. Ojha, and N. Spooner. Fractal: Post-quantum and transparent recursive proofs from holography. In *Advances in Cryptology–EUROCRYPT 2020: 39th Annual International Conference on the Theory and Applications of Cryptographic Techniques, Zagreb, Croatia, May 10–14, 2020, Proceedings, Part I 39*, pages 769–793. Springer, 2020.

CS10.   C. Costello and D. Stebila. Fixed argument pairings. In *Progress in Cryptology–LATINCRYPT 2010: First International Conference on Cryptology and Information Security in Latin America, Puebla, Mexico, August 8-11, 2010, proceedings 1*, pages 92–108. Springer, 2010.

D+15.   M. J. Dworkin et al. Sha-3 standard: Permutation-based hash and extendable-output functions. 2015.

Dev24.   Z. Developers. Halo2 proofs documentation. https://zcash.github.io/halo2/, 2024. Accessed: 2024-04-07.

Die11.   C. Diem. On the discrete logarithm problem in elliptic curves. *Compositio Mathematica*, 147(1):75–104, 2011.

dMU.   L. de Moura and S. Ullrich. The lean 4 theorem prover and programming language (system description).

FS87.   A. Fiat and A. Shamir. How to prove yourself: Practical solutions to identification and signature problems. In *Advances in Cryptology — CRYPTO' 86*, pages 186–194, Berlin, Heidelberg, 1987. Springer Berlin Heidelberg.

GGPR13.   R. Gennaro, C. Gentry, B. Parno, and M. Raykova. Quadratic span programs and succinct nizks without pcps. In *Advances in Cryptology–EUROCRYPT 2013: 32nd Annual International Conference on the Theory and Applications of Cryptographic Techniques, Athens, Greece, May 26-30, 2013. Proceedings 32*, pages 626–645. Springer, 2013.

GKM+18.   J. Groth, M. Kohlweiss, M. Maller, S. Meiklejohn, and I. Miers. Updatable and universal common reference strings with applications to zk-snarks. In *Annual International Cryptology Conference*, pages 698–728. Springer, 2018.

GKR+21.   L. Grassi, D. Khovratovich, C. Rechberger, A. Roy, and M. Schofnegger. Poseidon: A new hash function for {Zero-Knowledge} proof systems. In *30th USENIX Security Symposium (USENIX Security 21)*, pages 519–535, 2021.

GMR19.   S. Goldwasser, S. Micali, and C. Rackoff. The knowledge complexity of interactive proof-systems. In *Providing sound foundations for cryptography: On the work of shafi goldwasser and silvio micali*, pages 203–225. 2019.

Gro16.   J. Groth. On the size of pairing-based non-interactive arguments. In *Advances in Cryptology–EUROCRYPT 2016: 35th Annual International Conference on the Theory and Applications of Cryptographic Techniques, Vienna, Austria, May 8-12, 2016, Proceedings, Part II 35*, pages 305–326. Springer, 2016.

GW11.   C. Gentry and D. Wichs. Separating succinct non-interactive arguments from all falsifiable assumptions. In *Proceedings of the forty-third annual ACM symposium on Theory of computing*, pages 99–108, 2011.

GW20a.   A. Gabizon and Z. J. Williamson. plookup: A simplified polynomial protocol for lookup tables. *Cryptology ePrint Archive*, 2020.

GW20b.   A. Gabizon and Z. J. Williamson. Proposal: The turbo-plonk program syntax for specifying snark programs, 2020.

GWC19.   A. Gabizon, Z. J. Williamson, and O. Ciobotaru. Plonk: Permutations over lagrange-bases for oecumenical noninteractive arguments of knowledge. *Cryptology ePrint Archive*, 2019.

Hab22.   U. Haböck. Multivariate lookups based on logarithmic derivatives. *Cryptology ePrint Archive*, 2022.

KDT24.   C. Kwan, Q. Dao, and J. Thaler. Verifying jolt zkvm lookup semantics. *Cryptology ePrint Archive*, 2024.

Kil92.   J. Kilian. A note on efficient zero-knowledge proofs and arguments. In *Proceedings of the twenty-fourth annual ACM symposium on Theory of computing*, pages 723–732, 1992.

KZG10.   A. Kate, G. M. Zaverucha, and I. Goldberg. Constant-size commitments to polynomials and their applications. In *Advances in Cryptology-ASIACRYPT 2010: 16th International Conference on the Theory and Application of Cryptology and Information Security, Singapore, December 5-9, 2010. Proceedings 16*, pages 177–194. Springer, 2010.

MBKM19.   M. Maller, S. Bowe, M. Kohlweiss, and S. Meiklejohn. Sonic: Zero-knowledge snarks from linear-size universal and updatable structured reference strings. In *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security*, pages 2111–2128, 2019.

Mic94.   S. Micali. Cs proofs. In *Proceedings 35th Annual Symposium on Foundations of Computer Science*, pages 436–453. IEEE, 1994.

Nak08.   S. Nakamoto. Bitcoin: A peer-to-peer electronic cash system. 2008.

OO25.   OpenVM-Org, Mar 2025.

PCW+23.   S. Pailoor, Y. Chen, F. Wang, C. Rodríguez, J. Van Geffen, J. Morton, M. Chu, B. Gu, Y. Feng, and I. Dillig. Automated detection of Under-Constrained circuits in Zero-Knowledge proofs. *Proc. ACM Program. Lang.*, 7(PLDI), June 2023.

PHGR16.   B. Parno, J. Howell, C. Gentry, and M. Raykova. Pinocchio: Nearly practical verifiable computation. *Communications of the ACM*, 59(2):103–112, 2016.

POL22.   POLYGON. Plonky2: Fast recursive arguments with plonk and fri. `https://github.com/0xPolygonZero/plonky2/blob/136cdd053f2175134cddc61abc587f1862e76921/plonky2/plonky2.pdf`, 2022. Accessed: 2025-05-10.

Pol24.   Polygon Labs. zkevm and execution traces- polygon knowledge layer, 2024.

PSS19.   A. Pertsev, R. Semenov, and R. Storm. Tornado cash privacy solution. Technical report, Tornado Cash, 2019.

RIS23.   RISC Zero. Proof system in detail. Technical report, RISC Zero, 2023. [Accessed 08-05-2025].

Scr24.   Scroll. zkevm overview, 2024.

Set20.   S. Setty. Spartan: Efficient and general-purpose zksnarks without trusted setup. In *Annual International Cryptology Conference*, pages 704–737. Springer, 2020.

Suc23.   Succinct. Introducing SP1, 2023.

WTS⁺18.   R. S. Wahby, I. Tzialla, A. Shelat, J. Thaler, and M. Walfish. Doubly-efficient zksnarks without trusted setup. In *2018 IEEE Symposium on Security and Privacy (SP)*, pages 926–943. IEEE, 2018.